1 NLP reproduce ability

1.1 Lora

Here is a reproduction of Lora tasks

https://github.com/microsoft/LoRA

The best accuracy gets at 13 epochs. However, it has a set of about 80epochs. At the 13th epoch, the best accuracy occurs.

Then it slowly decays.

Lora gives a reproduction of SOTA by using huggingface tutorial codes. But it only trains 3 epochs, which is unacceptable.

1.2 Fairseq

Also, fairest provides a set of parameters.

https://github.com/facebookresearch/fairseq/blob/main/examples/roberta/config/finetuning/cola.yaml

It only trains10 epochs. But the same problems occurs.

fairseq settings COLA finetune

And my settings generate this curve.

Settings

Backend	Roberta-base
Dataset	COLA
Batch size	32
Lr	2e-5
Weight decay	1e-4
Optimizer	Adam
Epochs	20

2 Pipeline training on two machines

2.1 Settings

Here I use 40CPUs to simulate the client and one GTX1080 to simulate the server

Backend	MobileNetV2
Dataset	CIFAR10
Batch size	64
Image size	[3,224,224]
Lr	0.005
Weight decay	0
Optimizer	SGD with momentum
Momentum	0.1
Partition	Client: Conv+bn, Classifier Server: The rest
Chunk	8

2.2 Results

Here are the definitions of the bandwidth.

$$Bandwidth_{avg} = data(send\ or\ recv)/calculation\_time$$
$$Bandwidth_{peak,client} = max(Bandwidth_{send})$$
$$Bandwidth_{peak,server} = max(Bandwidth_{recv})$$

And there are two parameters for power iteration.

As we all know. power iteration is a way of PCA which has similar to SVD_lowrank. But SVD is a high-cost algorithm. Power iteration cost less since it uses QR decomposition instead of SVD decomposition. But QR decomposition could cost a lot when the rank is bigger.

For CV tasks. Activation memory has a size of `[B,C,H,W]`. I unsqueeze the last two ranks to `[B,C,H*W]`. And then I use power iteration to spread it to `[B,C,rank]` and `[B,rank,H*W]`. As you can see, power iteration(3,7), means two ranks of the two sizes of activation memory.

Compress Method	Compress Rate	Acc	Bandwidth(Avg)	Bandwidth(Peak,Client)	Bandwidth(Peak,Server)	Computation Time	Total Time per batch	Chunk
None	1.0	95.86	236.07MB/s	455MB/s	454MB/s	0.48s	2.47s	1
None	1.0	96.04	343.37MB/s	460MB/s	457MB/s	0.33s	2.15s	4
None	1.0	95.94	323.75MB/s	454MB/s	458MB/s	0.35s	2.07s	8
Conv Insert	0.097	96.01	18.10MB/s	147MB/s	149MB/s	0.55s	0.38s	1
Conv Insert	0.097	96.01	29.27MB/s	251MB/s	253MB/s	0.34s	0.38s	4
Conv Insert	0.097	96.01	27.65MB/s	252MB/s	253MB/s	0.36s	0.38s	8
Conv Insert	0.070	95.89	13.92MB/s	154MB/s	148MB/s	0.55s	0.62s	1
Conv Insert	0.070	95.89	22.52MB/s	254MB/s	255MB/s	0.34s	0.37s	4
Conv Insert	0.070	95.87	21.27MB/s	258MB/s	261MB/s	0.36s	0.37s	8
Poweriter(3,7)	0.101	95.54	20.50MB/s	174MB/s	177MB/s	0.56s	0.66s	1
Poweriter(3,7)	0.101	95.54	32.80MB/s	269MB/s	249MB/s	0.35s	0.43s	4
Poweriter(3,7)	0.101	95.54	31.03MB/s	246MB/s	248MB/s	0.37s	0.43s	4