

Weekly Report 6.6

Kanghao

June 7, 2022

1 NLP reproduce

1.1 LoRA

Here is a reproduction of [LORA](#).

The best accuracy gets at 13 epochs. However, it has a set of about 80epochs. At the 13th epoch, the best accuracy occurs.

Then it slowly decays.

There is a graph about LORA implemented Roberta-base finetune COLA dataset, the y-axis is validation matthews correlation, and the x-axis is epoch number.

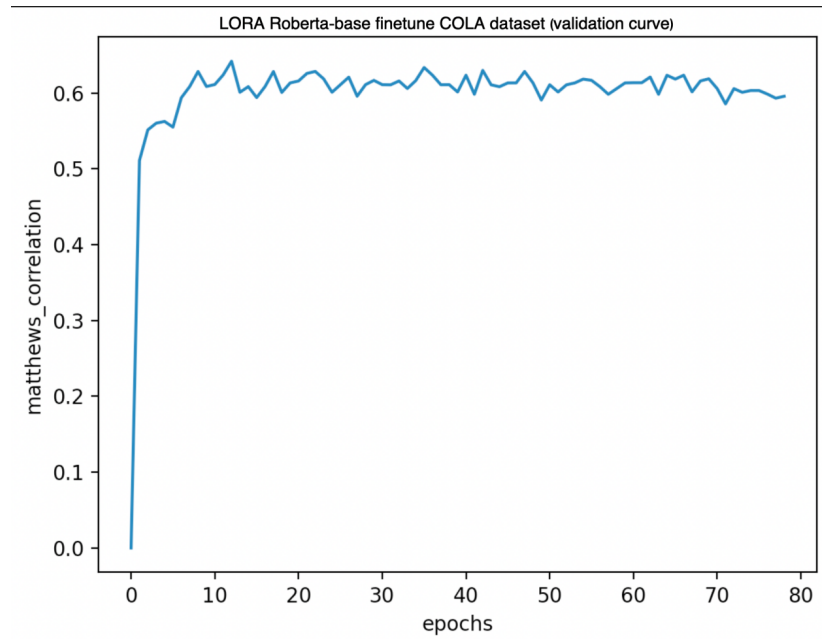


Figure 1: LORA Roberta-base finetune on COLA datasets

1.2 Fairseq

Also, fairseq provides a set of [parameters](#).

Backend	Roberta-base
Dataset	COLA
Batch size	16
Lr	1e-5
Weight decay	0.1
Optimizer	Adam(eps = 1e-6)
Epochs	10
Method	mixed precision

It only trains 10 epochs. But the same problems occurs. The best validation accuracy occurs at the third epoch. Here is a curve showing overfitting problems.

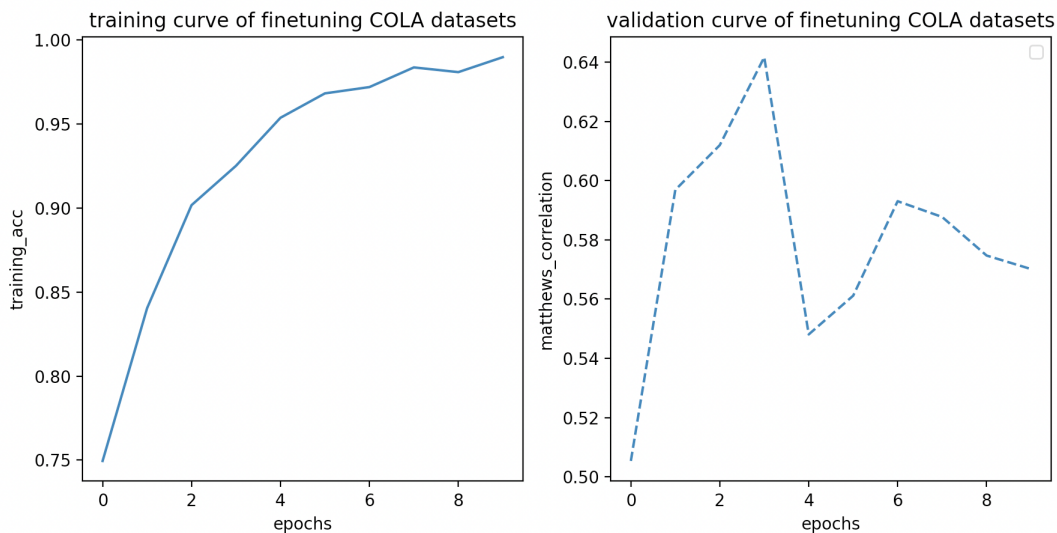


Figure 2: Fairseq Roberta-base finetune COLA datasets

1.3 Gpipe settings

Backend	Roberta-base
Dataset	COLA
Batch size	32
Lr	2e-5
Weight decay	1e-4
Optimizer	Adam(eps = 1e-6)
Epochs	20
Method	mixed precision

And my settings generate this curve.

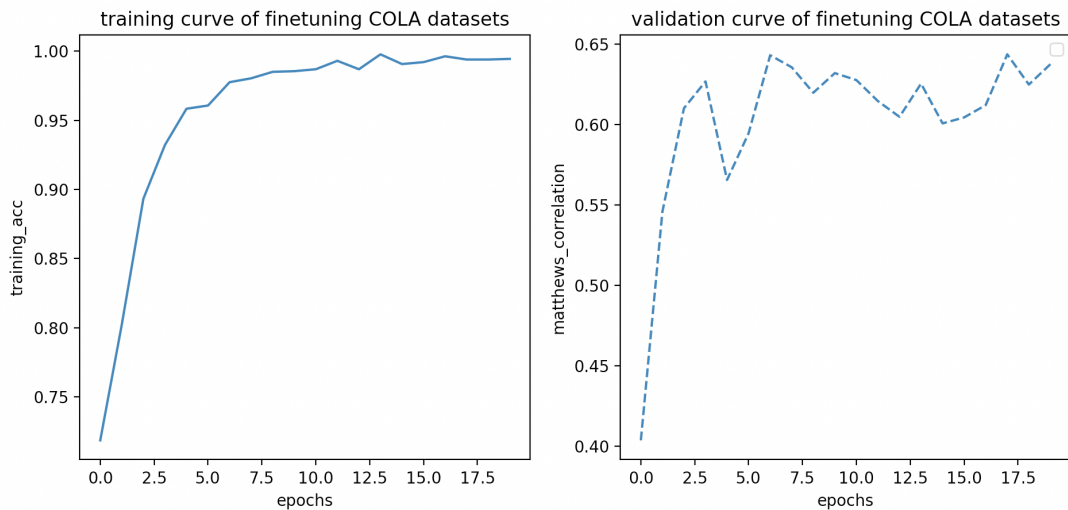


Figure 3: Dataparallel Roberta-base finetune COLA datasets(20 epochs) with same settings shown on the table

If we change epoch to 40, the curves are shown below.

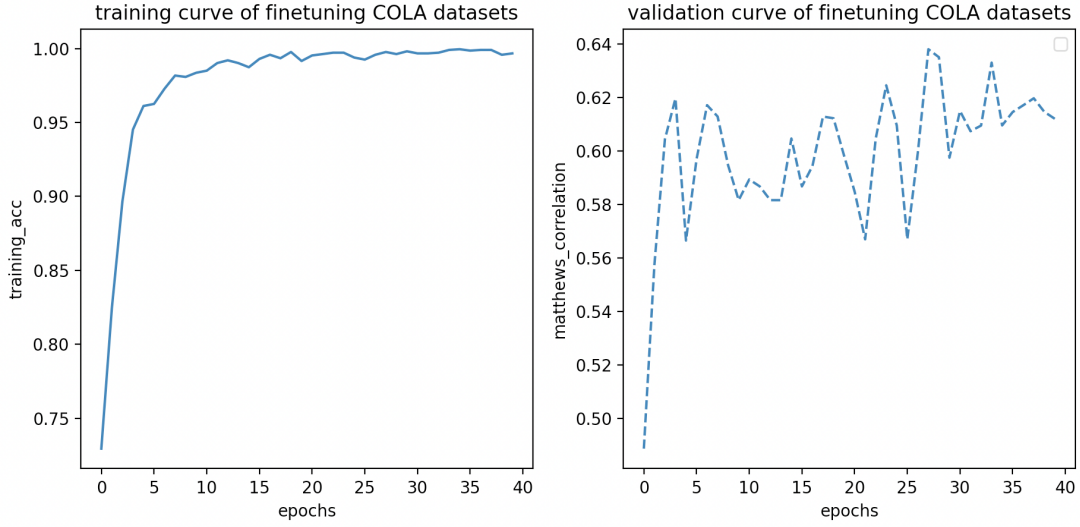


Figure 4: Dataparallel Roberta-base finetune COLA datasets (40epochs) with same settings shown on the table

Also when finetuning COLA dataset with Power iteration of 16ranks. The settings gets an acceptable result.

Backend	Roberta-base
Dataset	COLA
Batch size	32
Lr	2e-5
Weight decay	1e-4
Optimizer	Adam(eps = 1e-6)
Epochs	20
Method	mixed precision
Partition	Client: self attention+dense+classifier,Server: fastfeedforward+robertalayer[1:11]

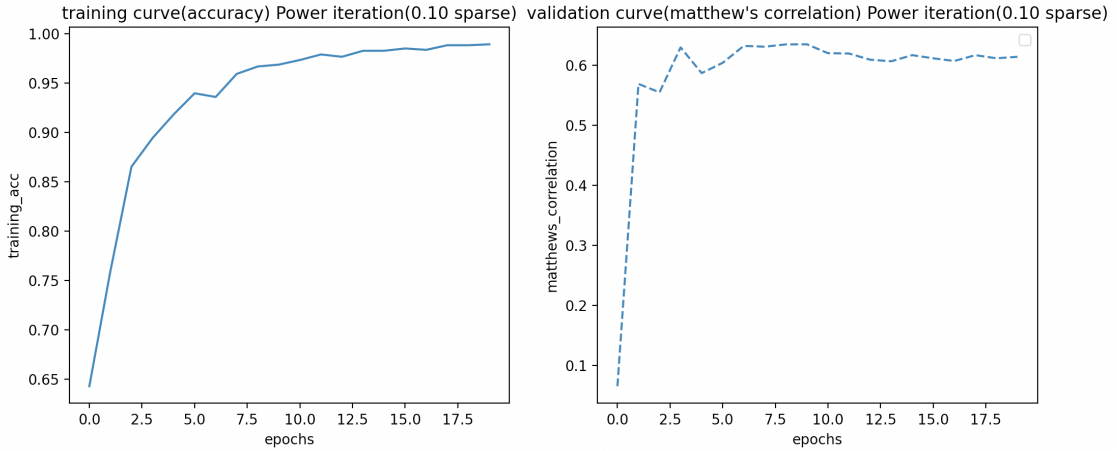


Figure 5: Using pipeline parallelism and settings shown in the table finetuning COLA dataset with Roberta-base and Power iteration of 16ranks

2 Pipeline training on two machines

2.1 Settings

Here I use 40CPUs to simulate the client and one GTX1080 to simulate the server.

Backend	MobileNetV2
Dataset	CIFAR10
Batch size	64
Image size	[3,224,224]
Lr	0.005
Weight decay	0.0
Optimizer	SGD with momentum
Momentum	0.1
Partition	Client: Conv+bn Classifier Server: Relu + features[1:]
Chunk	1,4,8

Also, when finetune with poweriteration

2.2 Results

Here are the definitions of the bandwidth.

$$Bandwidth_{avg} = data(send\ or\ recv)/computation_time$$

$$Bandwidth_{peak,client} = max(Bandwidth_{send})$$

$$Bandwidth_{peak,server} = max(Bandwidth_{recv})$$

And there are two parameters for power iteration.

As we all know. power iteration is a way of PCA which has similar to SVD_lowrank. But SVD is a high-cost algorithm. Power iteration cost less since it uses QR decomposition instead of SVD decomposition. But QR decomposition could cost a lot when the rank is bigger.

For CV tasks. Activation memory has a size of [B,C,H,W]. I unsqueeze the last two ranks to [B,C,H*W]. And then I use power iteration to spread it to [B,C,rank] and [B,rank,H*W].

As you can see, power iteration(3,7), means two ranks of the two sizes of activation memory. Activation memory with size [64,32,112,112] is compressed to [64,3,112*112] and [64,32,3]. Activation memory with size [64,1280,7,7] is compressed to [64,7,7*7] and [64,1280,7].

For conv insert, I use convolution and transpose convolution to compress activation memory. Conv Insert with compress rate 0.097 uses convolution2d with 32 in_channels, 32 out_channels and (4,4) kernel to compress the first set of activation memory([64,32,112,112]), and it uses convolution2d with 1280 in_channels and 320 out_channels to compress the last set of activation memory([64,1280,7,7]). Conv Insert with compress rate 0.070 uses convolution2d with 32 in_channels, 20 out_channels and (4,4) kernel to compress the first set of activation memory([64,32,112,112]), and it uses convolution2d with 1280 in_channels and 320 out_channels to compress the last set of activation memory([64,1280,7,7]).

Compress Method	Compress Rate	Acc	Bandwidth(Avg)	Bandwidth(Peak,Client)	Bandwidth(Peak,Server)	Computation Time	Total Time per batch	Chunk
None	1.0	95.86	236.07MB/s	455MB/s	454MB/s	0.48s	2.47s	1
None	1.0	96.04	343.37MB/s	460MB/s	457MB/s	0.33s	2.15s	4
None	1.0	95.94	323.75MB/s	454MB/s	458MB/s	0.35s	2.07s	8
Conv Insert	0.097	96.01	18.10MB/s	147MB/s	149MB/s	0.55s	0.38s	1
Conv Insert	0.097	96.01	29.27MB/s	251MB/s	253MB/s	0.34s	0.38s	4
Conv Insert	0.097	96.01	27.65MB/s	252MB/s	253MB/s	0.36s	0.38s	8
Conv Insert	0.070	95.89	13.92MB/s	154MB/s	148MB/s	0.55s	0.62s	1
Conv Insert	0.070	95.89	22.52MB/s	254MB/s	255MB/s	0.34s	0.37s	4
Conv Insert	0.070	95.87	21.27MB/s	258MB/s	261MB/s	0.36s	0.37s	8
Poweriter(3,7)	0.101	95.54	20.50MB/s	174MB/s	177MB/s	0.56s	0.66s	1
Poweriter(3,7)	0.101	95.54	32.80MB/s	269MB/s	249MB/s	0.35s	0.43s	4
Poweriter(3,7)	0.101	95.54	31.03MB/s	246MB/s	248MB/s	0.37s	0.43s	4

Here are visualization works. Peak bandwidth with different chunks and different compress algorithms.

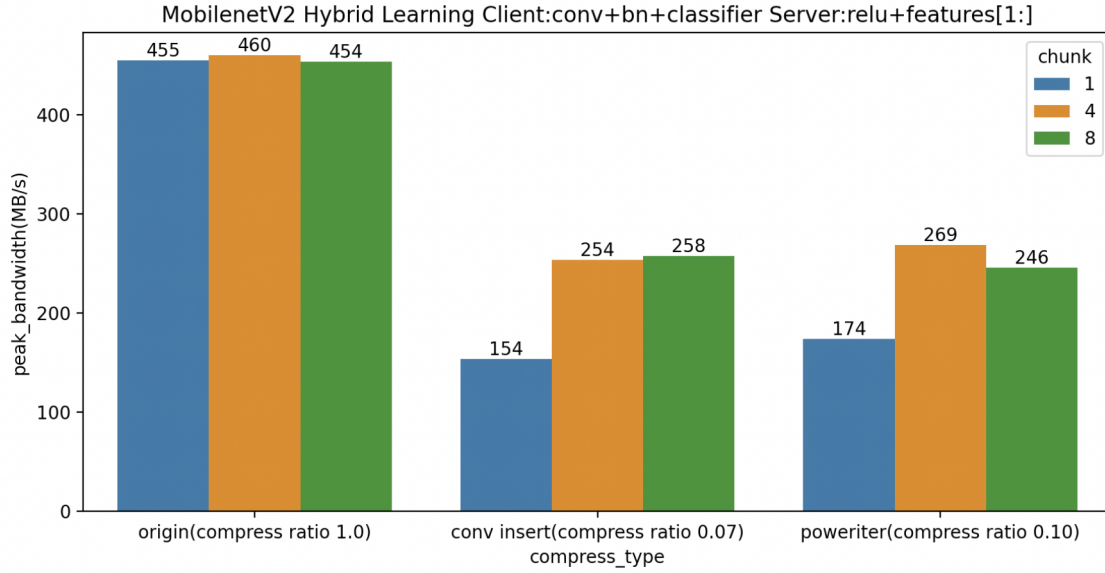


Figure 6: Peak bandwidth with different chunks and compress algorithms at two different machines using Pipeline parallelism finetune MobilenetV2 with CIFAR10

Average bandwidth with different chunks and different compress algorithms.

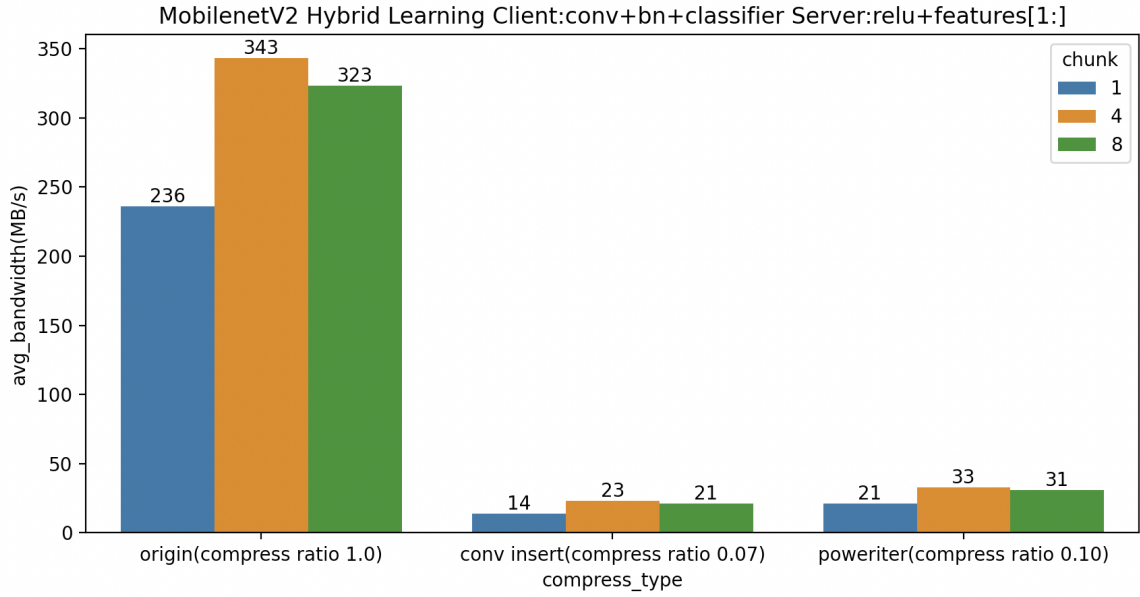


Figure 7: Average bandwidth with different chunks and compress algorithms at two different machines using Pipeline parallelism finetune MobilenetV2 with CIFAR10

Time per batch with different chunks and different compress algorithms.

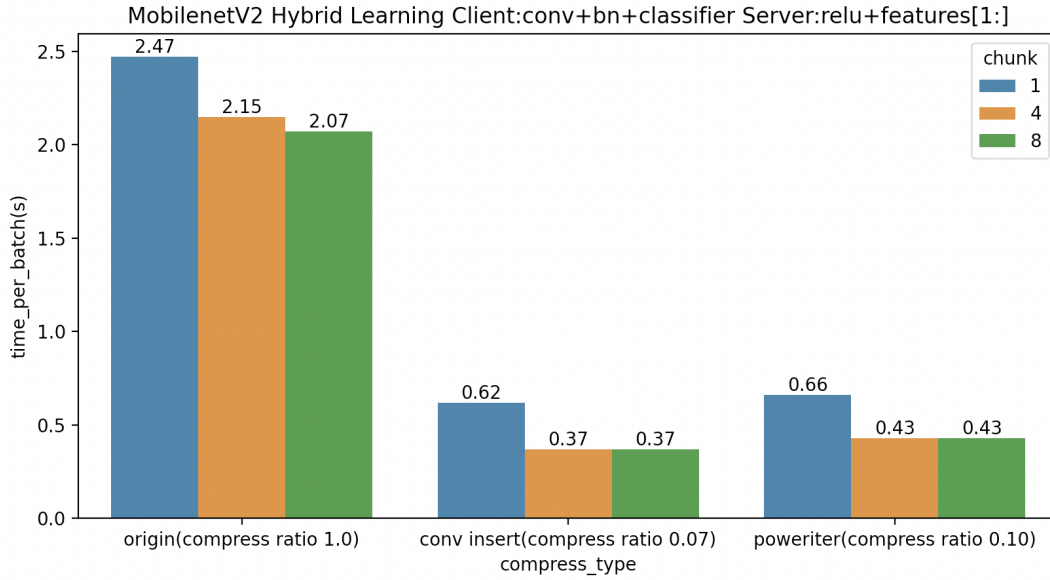


Figure 8: Time per batch with different chunks and different compress algorithms at two different machines using Pipeline parallelism finetune MobilenetV2 with CIFAR10

Computation time per batch with different chunks and different compress algorithms.

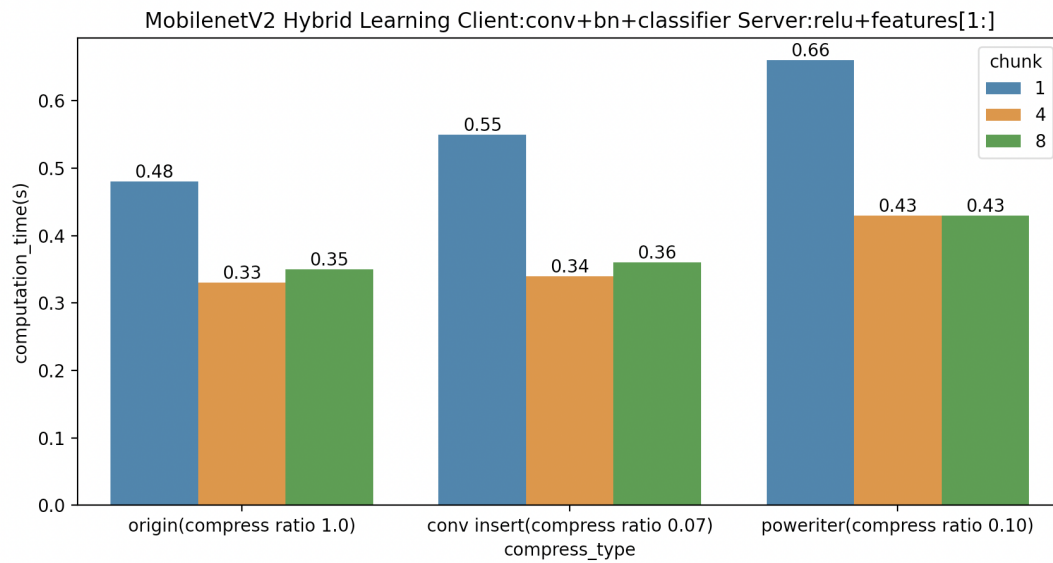


Figure 9: Computation time per batch with different chunks and different compress algorithms at two different machines using Pipeline parallelism finetune MobilenetV2 with CIFAR10